

Computational Tools for Macroeconometrics

Assignment 3

This assignment covers two topics: Monte Carlo simulations and the Bootstrap with a focus on time series models. As usual, I will cover the basic ideas behind the concepts and then present the assignment.

Monte Carlo experiments

Monte Carlo experiments is a handy tool to assess the quality of the asymptotic approximation of econometric estimators.

Consider the following linear model

$$y_t = \beta_0^o + \beta_1^o x_t + u_t, t = 1, \dots, T,$$

where y_t , x_t , and u_t are random variables and β_0^o and β_1^o are parameters to be estimated. In matrix form, the model can be written as

$$Y_t = \mathbf{X}_t \beta^o + U_t,$$

where

$$\underline{Y}_{(T \times 1)} = \begin{pmatrix} Y_1 \\ \vdots \\ Y_T \end{pmatrix}, \underline{X}_{T \times 2} = \begin{pmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_T \end{pmatrix}, \underline{U}_{(T \times 1)} = \begin{pmatrix} u_1 \\ \vdots \\ u_T \end{pmatrix}, \underline{\beta^o}_{(2 \times 1)} = \begin{pmatrix} \beta_0^o \\ \beta_1^o \end{pmatrix}$$

We will assume that 1. (u_t, x_t) are i.i.d. over t ; 2. $E(u_t | x_t) = 0$, $t = 1, \dots, T$; 3. $E(|x_t|^4) < \infty$ and $E(|u_t|^4) < \infty$; 4. $E(u_t^2 | x_t) = \sigma^2 > 0$.

Under these assumptions, the OLS estimator of β^o ,

$$\hat{\beta}^{ols} = (\mathbf{X}'\mathbf{X})^{-1} \mathbf{X}'Y$$

is unbiased, consistent, and asymptotically normal:

$$E(\hat{\beta}^{ols}) = \beta^o, \quad \hat{\beta}^{ols} \xrightarrow{p} \beta^o, \quad (\hat{\beta}^{ols} - \beta^o) \xrightarrow{d} N[0, \sigma^2 E(\mathbf{X}'\mathbf{X})^{-1}].$$

These results are derived theoretically. For unbiasedness, the law of iterated expectations and the iid of (u_t, x_t) gives

$$E(E(\hat{\beta}^{ols} | \mathbf{X})) = \beta^o + E((\mathbf{X}'\mathbf{X})^{-1} \mathbf{X}' E(U | \mathbf{X})) = \beta^o.$$

Under the iid assumption and the restrictions on the fourth moments of u_t and x_t , we have

$$\frac{1}{T} \mathbf{X}'U = \frac{1}{n} \sum_{t=1}^T \begin{pmatrix} u_t \\ x_t u_t \end{pmatrix} \xrightarrow{p} \begin{pmatrix} 0 \\ 0 \end{pmatrix}, \quad \text{and} \quad \frac{1}{T} \mathbf{X}'\mathbf{X} = \frac{1}{T} \sum_{t=1}^T \begin{pmatrix} 1 & x_t \\ x_t & x_t^2 \end{pmatrix} \xrightarrow{p} \begin{pmatrix} 1 & E(x_t) \\ E(x_t) & E(x_t^2) \end{pmatrix}.$$

These convergences in probability deliver consistency of the OLS estimator

$$\hat{\beta}^{ols} = \beta^o + \left(\frac{1}{T} \mathbf{X}' \mathbf{X} \right)^{-1} \frac{1}{T} \mathbf{X}' U \xrightarrow{p} \beta^o + \begin{pmatrix} 1 & E(x_t) \\ E(x_t) & E(x_t^2) \end{pmatrix}^{-1} \begin{pmatrix} 0 \\ 0 \end{pmatrix} = \beta^o.$$

Finally, iid, the moments' restrictions, and the homoskedasticity assumptions, give

$$\sqrt{T} (\hat{\beta}^{ols} - \beta^o) \xrightarrow{d} N \left[0, \sigma^2 \begin{pmatrix} 1 & E(x_t) \\ E(x_t) & E(x_t^2) \end{pmatrix}^{-1} \right].$$

Consistency and asymptotic results, and they hold as $T \rightarrow \infty$. Given a sample of size T , say $T = 50$, how close is the distribution of $\hat{\beta}^o$ the normal one postulated by the usual asymptotic theory? How close is $\hat{\beta}^o$ to β^o ? Unfortunately, theory only tells us that the larger the sample size, the better the normal approximation. Similarly, it tells us that $\hat{\beta}^o - \beta^o$ is smaller the larger the sample size.

Monte Carlo simulations can help us determine how good these approximations are for a given size of T . The idea is simple: we simulate a sample of (u_t, x_t) , $t = 1, \dots, T$ from which we generate y_t (using a arbitrary value for β_0^o and β_1^o .) We estimate the parameters using the simulated data. We repeat this operation many times, saving the estimated parameters. The saved parameters give the empirical distribution of the OLS estimator and help verify how closely the theory matches the empirical distribution.

The following code does a Monte Carlo for the linear model above:

```

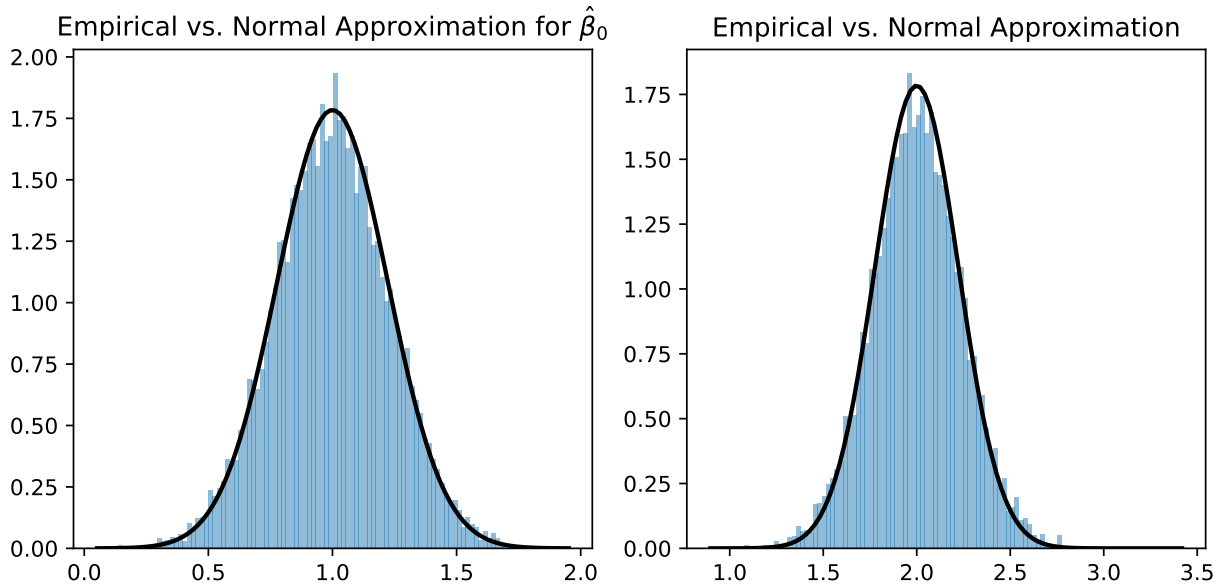
1  import numpy as np
2  import matplotlib.pyplot as plt
3  from scipy.stats import norm
4
5  # Set parameters
6  T = 20
7  beta_0_true = 1
8  beta_1_true = 2
9  sigma = 1
10 num_simulations = 10000
11
12 # Arrays to store the estimates from each simulation
13 beta_0_estimates = np.zeros(num_simulations)
14 beta_1_estimates = np.zeros(num_simulations)
15
16 # Run simulations
17 for i in range(num_simulations):
18     x = np.random.normal(0, 1, T)
19     u = np.random.normal(0, sigma, T)
20     y = beta_0_true + beta_1_true * x + u
21
22     # OLS estimation
23     X = np.vstack([np.ones(T), x]).T
24     beta_hat = np.linalg.inv(X.T @ X) @ X.T @ y
25     beta_0_estimates[i] = beta_hat[0]
26     beta_1_estimates[i] = beta_hat[1]
27
28 # Plotting the results
29 fig, ax = plt.subplots(1, 2, figsize=(8, 4))

```

```

30
31 # Distribution of beta_0
32 ax[0].hist(beta_0_estimates, bins=100, alpha=0.5, density=True)
33 xmin, xmax = ax[0].get_xlim()
34 x = np.linspace(xmin, xmax, 100)
35 p = norm.pdf(x, beta_0_true, 1/np.sqrt(T))
36 ax[0].plot(x, p, 'k', linewidth=2)
37 ax[0].set_title(f'Empirical vs. Normal Approximation for  $\hat{\beta}_0$ ')
38
39 # Distribution of beta_1
40 ax[1].hist(beta_1_estimates, bins=100, alpha=0.5, density=True)
41 xmin, xmax = ax[1].get_xlim()
42 x = np.linspace(xmin, xmax, 100)
43 p = norm.pdf(x, beta_1_true, 1/np.sqrt(T))
44 ax[1].plot(x, p, 'k', linewidth=2)
45 ax[1].set_title(f'Empirical vs. Normal Approximation')
46
47 plt.tight_layout()
48 plt.show()

```



The approximation is excellent, even if $T = 20$. This might sound surprising since the normal distribution of $\hat{\beta}^{ols}$ is normal when T is large. The reason why we get such a close agreement between the empirical distribution of $\hat{\beta}^{ols}$ and the theoretical one is that we are simulating u_t the data from a normal distribution independently from x_t . When $u_t|x_t \sim N(0, \sigma^2)$, the distribution of the OLS estimator is precisely normal, even when $T = 3$. Of course, if T is small, the estimator's variance will be larger, but the shape of the distribution of the OLS will be normal.

Instead of squinting at the histograms and the density implied by the CLT, we can modify the code to calculate the confidence interval at each simulation and see how many times the true values of the parameters fall into all the intervals generated. If the approximation is good, a 95% confidence interval should contain the true parameters 95% of the time.

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from scipy.stats import norm
4
5  # Set parameters
6  T = 50
7  beta_0_true = 1
8  beta_1_true = 2
9  sigma = 1
10 num_simulations = 10000
11
12 # Arrays to store the estimates from each simulation
13 beta_0_estimates = np.zeros(num_simulations)
14 beta_1_estimates = np.zeros(num_simulations)
15 beta_0_in = np.zeros(num_simulations)
16 beta_1_in = np.zeros(num_simulations)
17
18 # Run simulations
19 for i in range(num_simulations):
20     x = np.random.normal(0,1,T)
21     u = np.random.normal(0,sigma,T)
22     y = beta_0_true + beta_1_true * x + u
23     # OLS estimation
24     X = np.vstack([np.ones(T), x]).T
25     XXinv = np.linalg.inv(X.T @ X)
26     beta_hat = XXinv @ X.T @ y
27     beta_0_estimates[i] = beta_hat[0]
28     beta_1_estimates[i] = beta_hat[1]
29     u_hat = y - beta_hat[0] - beta_hat[1] * x
30     sigma2_hat = np.dot(u_hat, u_hat)/(T-2)
31     variance_hat = sigma2_hat*XXinv
32     se_0 = np.sqrt(variance_hat[0,0])
33     se_1 = np.sqrt(variance_hat[1,1])
34     ## Check whether beta_0 in CI 95%
35     beta_0_in[i] = beta_hat[0] - 1.965*se_0 < beta_0_true < beta_hat[0] + 1.965*se_0
36     beta_1_in[i] = beta_hat[1] - 1.965*se_1 < beta_1_true < beta_hat[1] + 1.965*se_1
37
38 # Output the results
39 print(f"Empirical 95% CI for beta_0: {np.mean(beta_0_in)}")
40 print(f"Empirical 95% CI for beta_1: {np.mean(beta_1_in)}")

```

Empirical 95% CI for beta_0: 0.9482

Empirical 95% CI for beta_1: 0.9497

The empirical coverage of the confidence intervals is only in the neighborhood of 95%. The reason is that we are estimating the variance of the coefficient, and in this case, the distribution of $\hat{\beta}_1/se(\hat{\beta}_1)$ has a t -student distribution with $T - 1$ degrees of freedom. If you run the code above for $T = 50$, the confidence interval coverage will be closer to 95%.

Let us see what happens if we simulate x_t and u_t from a chi-squared distribution with 5 of freedom centered in a way to have mean zero and unit variance.

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from scipy.stats import norm
4
5  # Set parameters
6  T = 50
7  beta_0_true = 1
8  beta_1_true = 2
9  sigma = 1
10 num_simulations = 10000
11
12 # Arrays to store the estimates from each simulation
13 beta_0_estimates = np.zeros(num_simulations)
14 beta_1_estimates = np.zeros(num_simulations)
15 beta_0_in = np.zeros(num_simulations)
16 beta_1_in = np.zeros(num_simulations)
17
18 # Run simulations
19 for i in range(num_simulations):
20     x = (np.random.chisquare(4,T) - 4)/np.sqrt(2*4)
21     u = (np.random.chisquare(4,T) - 4)/np.sqrt(2*4)
22     y = beta_0_true + beta_1_true * x + u
23     # OLS estimation
24     X = np.vstack([np.ones(T), x]).T
25     XXinv = np.linalg.inv(X.T @ X)
26     beta_hat = XXinv @ X.T @ y
27     beta_0_estimates[i] = beta_hat[0]
28     beta_1_estimates[i] = beta_hat[1]
29     u_hat = y - beta_hat[0] - beta_hat[1] * x
30     sigma2_hat = np.dot(u_hat, u_hat)/(T-2)
31     variance_hat = sigma2_hat*XXinv
32     se_0 = np.sqrt(variance_hat[0,0])
33     se_1 = np.sqrt(variance_hat[1,1])
34     ## Check whether beta_0 in CI 95%
35     beta_0_in[i] = beta_hat[0] - 1.965*se_0 < beta_0_true < beta_hat[0] + 1.965*se_0
36     beta_1_in[i] = beta_hat[1] - 1.965*se_1 < beta_1_true < beta_hat[1] + 1.965*se_1
37
38 # Output the results
39 print(f"Empirical 95% CI for beta_0: {np.mean(beta_0_in)}")
40 print(f"Empirical 95% CI for beta_1: {np.mean(beta_1_in)}")

```

Empirical 95% CI for beta_0: 0.9351
Empirical 95% CI for beta_1: 0.9463

Even a $T = 50$, there is some difference between the empirical and theoretical coverage.

We can use the simulations to see what happens when the conditions on the moments of u_t and x_t are unsatisfied.

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from scipy.stats import norm
4
5  # Set parameters
6  T = 150
7  beta_0_true = 1
8  beta_1_true = 2
9  sigma = 1
10 num_simulations = 10000
11
12 # Arrays to store the estimates from each simulation
13 beta_0_estimates = np.zeros(num_simulations)
14 beta_1_estimates = np.zeros(num_simulations)
15 beta_0_in = np.zeros(num_simulations)
16 beta_1_in = np.zeros(num_simulations)
17
18 # Run simulations
19 for i in range(num_simulations):
20     x = np.random.standard_cauchy(T)
21     u = np.random.standard_cauchy(T)
22     y = beta_0_true + beta_1_true * x + u
23     # OLS estimation
24     X = np.vstack([np.ones(T), x]).T
25     XXinv = np.linalg.inv(X.T @ X)
26     beta_hat = XXinv @ X.T @ y
27     beta_0_estimates[i] = beta_hat[0]
28     beta_1_estimates[i] = beta_hat[1]
29     u_hat = y - beta_hat[0] - beta_hat[1] * x
30     sigma2_hat = np.dot(u_hat, u_hat)/(T-2)
31     variance_hat = sigma2_hat*XXinv
32     se_0 = np.sqrt(variance_hat[0,0])
33     se_1 = np.sqrt(variance_hat[1,1])
34     ## Check weather beta_0 in CI 95%
35     beta_0_in[i] = beta_hat[0] - 1.965*se_0 < beta_0_true < beta_hat[0] + 1.965*se_0
36     beta_1_in[i] = beta_hat[1] - 1.965*se_1 < beta_1_true < beta_hat[1] + 1.965*se_1
37
38 # Output the results
39 print(f"Empirical 95% CI for beta_0: {np.mean(beta_0_in)}")
40 print(f"Empirical 95% CI for beta_1: {np.mean(beta_1_in)}")

```

Empirical 95% CI for beta_0: 0.9799
Empirical 95% CI for beta_1: 0.958

Bootstrap

The bootstrap is a statistical tool used for estimating the distribution of a statistic based on random sampling with replacement. It allows for assessing the variability of a statistic, and it is beneficial in settings where the theoretical distribution of the statistic is unknown or difficult to derive.

In the context of a linear regression model, the bootstrap can be used to perform inference on the regression coefficients. Here is how it works:

1. **Fit the Model:** We fit the linear regression model to the data to estimate the coefficients, say $\hat{\beta}$ and $\hat{\beta}_1$.
2. **Resample Data:** Generate many bootstrap samples by randomly sampling with replacement from the original dataset. Each bootstrap sample should be the same size as the original dataset.
- 3.
4. **Re-estimate Parameters:** For each bootstrap sample, refit the regression model to estimate the coefficients. This will yield a new set of estimates, $\hat{\beta}^*$.
5. **Calculate Statistics:** Calculate the variance of the bootstrap estimates.

The following Python code calculate the standard errors of the linear model

```
1  import numpy as np
2
3  # Set random seed for reproducibility
4  np.random.seed(0)
5
6  # Generate some sample data
7  T = 100 # Number of observations
8  x = np.random.normal(0, 1, T)
9  u = np.random.normal(0, 1, T)
10 beta_0_true = 1
11 beta_1_true = 2
12
13 # Simulate response variable y
14 y = beta_0_true + beta_1_true * x + u
15
16 # Function to fit linear model
17 def fit_linear_model(x, y):
18     X = np.vstack([np.ones(len(x)), x]).T
19     beta_hat = np.linalg.inv(X.T @ X) @ (X.T @ y)
20     return beta_hat
21
22 # Initial fit
23 initial_beta = fit_linear_model(x, y)
24
25 # Number of bootstrap samples
26 B = 1000
27 bootstrap_estimates = np.zeros((B, 2))
28
29 # Perform bootstrap resampling
30 for i in range(B):
31     indices = np.random.choice(range(T), size=T, replace=True)
32     x_resampled = x[indices]
33     y_resampled = y[indices]
34     bootstrap_estimates[i] = fit_linear_model(x_resampled, y_resampled)
35
```

```

36 # Compute standard errors
37 standard_errors = bootstrap_estimates.std(axis=0)
38
39 print("Bootstrap Standard Errors:")
40 print("SE(beta_0):", standard_errors[0])
41 print("SE(beta_1):", standard_errors[1])
42
43 print("LM Standard Errors:")
44 import statsmodels.api as sm
45 X = sm.add_constant(x)
46 model = sm.OLS(y, X)
47 results = model.fit()
48
49 # Standard errors from statsmodels
50 statsmodels_se = results.bse
51 print("Standard Errors from statsmodels OLS:")
52 print("SE(beta_0):", statsmodels_se[0])
53 print("SE(beta_1):", statsmodels_se[1])

```

```

Bootstrap Standard Errors:
SE(beta_0): 0.10061319113712466
SE(beta_1): 0.09469134082416596
LM Standard Errors:
Standard Errors from statsmodels OLS:
SE(beta_0): 0.10404543947505013
SE(beta_1): 0.10305046686003076

```

Linear model with dependent data

So far, we have assumed that (u_t, x_t) are iid. With macroeconomic data, it is often the case that (u_t, x_t) are correlated.

Consider the following model:

$$y_t = \beta_0 + \beta_1 x_t + u_t,$$

with

$$x_t = \phi_x x_{t-1} + \eta_t, \quad |\phi_x| < 1$$

and

$$u_t = \phi_u u_{t-1} + \varepsilon_t, \quad |\phi_u| < 1.$$

If η_i and ε_t are independent, we will have that $E(u_t | x_t) = 0$ and so we can consistently estimate β_1 by OLS. The asymptotic distribution of the OLS estimator will be normal, but the variance of this distribution is difficult to estimate. If we use the standard errors that do not take into account the correlation over time of the variables, the inference will be off. We perform a simple Monte Carlo to see the impact of serial correlation in x_t and u_t .

```

1 def simulate_ar1(n, phi, sigma):
2     """
3     Simulate an AR(1) process.

```



```

4
5     Parameters:
6         n (int): Number of observations.
7         phi (float): Coefficient of AR(1) process.
8         sigma (float): Standard deviation of the innovation term.
9
10    Returns:
11        np.array: Simulated AR(1) error terms.
12    """
13    errors = np.zeros(n)
14    eta = np.random.normal(0, sigma, n) # white noise
15    for t in range(1, n):
16        errors[t] = phi * errors[t - 1] + eta[t]
17    return errors
18
19 def simulate_regression_with_ar1_errors(n, beta0, beta1, phi_x, phi_u, sigma):
20     """
21     Simulate a regression model with AR(1) error terms.
22
23     Parameters:
24         n (int): Number of observations.
25         beta0 (float): Intercept of the regression model.
26         beta1 (float): Slope of the regression model.
27         phi (float): Coefficient of the AR(1) process in the error term.
28         sigma (float): Standard deviation of the innovation term in the AR(1) process.
29
30     Returns:
31         tuple: x (independent variable), y (dependent variable), errors (AR(1) process)
32     """
33     x = simulate_ar1(n, phi_x, sigma)
34     u = simulate_ar1(n, phi_u, sigma)
35     y = beta0 + beta1 * x + u
36     return x, y, u
37
38 T = 500                # Number of observations
39 beta0 = 1.             # Intercept
40 beta1 = 2              # Slope
41 phi_x = 0.7            # AR(1) coefficient for x
42 phi_u = 0.7            # AR(1) coefficient for the errors
43 sigma = 1              # Standard deviation of the white noise
44
45 # Simulating the model
46
47 ## Do monte carlo
48 t_stats_hc = []
49 t_stats_hac = []
50
51 for i in range(1000):
52     x, y, errors = simulate_regression_with_ar1_errors(T, beta0, beta1, phi_x, phi_u, sigma)
53     X = sm.add_constant(x)

```

```

54     model = sm.OLS(y, X).fit(cov_type='HC1')
55     t_stats_hc.append(model.t_test('x1=2').tvalue)
56     ## Use HAC: takes into account serial correlation
57     model2 = sm.OLS(y, X).fit(cov_type='HAC', cov_kwds={'maxlags': np.floor(1.3*T**(1/2)).astype(int)})
58     t_stats_hac.append(model2.t_test('x1=2').tvalue)
59
60     ## Check we reject the null hypothesis at alpha=0.05 about 5% of the time
61
62     print(f"Empirical size test beta_1=2 using White SE: {np.mean(np.abs(np.array(t_stats_hc)) > 1.965)}")
63     print(f"Empirical size test beta_1=2 using HAC SE: {np.mean(np.abs(np.array(t_stats_hac)) > 1.965)}")

```

Empirical size test beta_1=2 using White SE: 0.255
Empirical size test beta_1=2 using HAC SE: 0.075

We can use the bootstrap to obtain the standard errors of $\hat{\beta}_1$. Unfortunately, with time-dependent data, the bootstrap needs to be modified. Instead of sampling with replacement observations from y_t and x_t , we will sample blocks of length ℓ with replacement. By resampling blocks, we ensure that the correlation in the data is preserved.

The moving block bootstrap (MBB) adapts the traditional bootstrap method to handle data where observations are dependent, such as in time series analysis. This approach involves resampling consecutive observation blocks to preserve the data's internal structure and dependence.

Steps for Moving Block Bootstrap (MBB):

1. **Choose Block Length:** Determine the length of the blocks, ℓ , that will be resampled. This length should be chosen based on the correlation structure of the data: the block length should be large enough to capture the dependence within the data.
2. **Generate Blocks:** From the original dataset of size T , generate new datasets by sampling blocks of length ℓ and concatenate them until reaching the size T . Blocks should be sampled with replacement.
3. **Resample Data within Blocks:** Sample y_t and x_t using the sampled blocks. This is the bootstrapped dataset.
4. **Refit the Model:** Fit the regression model to each bootstrapped dataset. Collect the estimated parameters for each bootstrapped dataset.
5. **Calculate Statistics:** Calculate the standard deviation of the bootstrap estimate. These are the standard errors.

```

1  import numpy as np
2  import statsmodels.api as sm
3
4  def moving_block_bootstrap(x, y, block_length, num_bootstrap):
5      T = len(y) # Total number of observations
6      num_blocks = T // block_length + (1 if T % block_length else 0)
7
8      # Fit the original model
9      X = sm.add_constant(x)
10     original_model = sm.OLS(y, X)

```

```

11     original_results = original_model.fit()
12
13     bootstrap_estimates = np.zeros((num_bootstrap, 2)) # Storing estimates for beta_0 and beta_1
14
15     # Perform the bootstrap
16     for i in range(num_bootstrap):
17         # Create bootstrap sample
18         bootstrap_indices = np.random.choice(np.arange(num_blocks) * block_length, size=num_blocks, repl
19         bootstrap_sample_indices = np.hstack([np.arange(index, min(index + block_length, T)) for index i
20         bootstrap_sample_indices = bootstrap_sample_indices[:T] # Ensure the bootstrap sample is the sa
21
22         x_bootstrap = x[bootstrap_sample_indices]
23         y_bootstrap = y[bootstrap_sample_indices]
24
25         # Refit the model on bootstrap sample
26         X_bootstrap = sm.add_constant(x_bootstrap)
27         bootstrap_model = sm.OLS(y_bootstrap, X_bootstrap)
28         bootstrap_results = bootstrap_model.fit()
29
30         # Store the estimates
31         bootstrap_estimates[i, :] = bootstrap_results.params
32
33     return bootstrap_estimates
34
35 # Run moving block bootstrap
36 block_length = 12
37 num_bootstrap = 1000
38 x, y, errors = simulate_regression_with_ar1_errors(200, beta0, beta1, phi_x, phi_u, sigma)
39 bootstrap_results = moving_block_bootstrap(x, y, block_length, num_bootstrap)
40
41 # Calculate and print standard errors
42 bootstrap_standard_errors = bootstrap_results.std(axis=0)
43 print("Bootstrap Standard Errors:")
44 print("SE(beta_0):", bootstrap_standard_errors[0])
45 print("SE(beta_1):", bootstrap_standard_errors[1])

```

```

Bootstrap Standard Errors:
SE(beta_0): 0.24226545211764944
SE(beta_1): 0.14191555909924308

```

Assignment 3

Task

Apply Monte Carlo simulations combined with bootstrap methods to evaluate the quality of inference on β_1 using serially correlated data.

Steps

1. Simulate data using `simulate_regression_with_ar1_errors` with parameters as specified in the document.
2. Calculate bootstrap standard errors.
3. Construct a 95% confidence interval for β_1 .
4. Perform Monte Carlo simulations for $T=100$ and $T=500$, and assess the empirical coverage of the confidence interval.

Assignment link: <https://classroom.github.com/a/GHyuRNbb>